LOVELY PROFESSIONAL UNIVERSITY

Sample Question Set - Unit 5

Course Code: CSE276
Course Title: ARTIFICIAL INTELLIGENCE FOUNDATIONS
Course Outcome: CO5
Paper: Sample Multiple Choice Questions
Maximum Marks: 40 x 1 = 40

Unit: 5
Unit Title: Generative AI, Responsible and Explainable AI
Unit Mapped RBT: L2 (Understand)
Total Questions: 40
Suggested Duration: 60 Minutes

Answer Grid (to be filled by the student)

Q. No.	Answer	Q. No.	Answer	Q. No.	Answer	Q. No.	Answer
Q1		Q11		Q21		Q31	
Q2		Q12		Q22		Q32	
Q3		Q13		Q23		Q33	
Q4		Q14		Q24		Q34	
Q5		Q15		Q25		Q35	
Q6		Q16		Q26		Q36	
Q7		Q17		Q27		Q37	
Q8		Q18		Q28		Q38	
Q9		Q19		Q29		Q39	
Q10		Q20		Q30		Q40	

Multiple Choice Questions (40 x 1 = 40 Marks)

Q1. Name the class of model that produces new content rather than assigning a label to an input.

A) A generative model
B) A discriminative model
C) A production system
D) A search algorithm

Q2. Identify the limit on how much text a large language model can take into account at one time.

A) The temperature setting
B) The context window
C) The token vocabulary
D) The parameter count

Q3. State what the role instruction of a prompt supplies to the model.

A) The maximum length of the reply
B) The documents it should search
C) The character or expertise it should adopt
D) The temperature it should use

Q4. Name the prompting technique that includes a few worked examples inside the prompt.

A) Zero-shot prompting
B) Chain-of-thought prompting
C) Retrieval grounding
D) Few-shot prompting

Q5. Identify the first step in preparing a document collection for Retrieval-Augmented Generation.

A) Fine-tuning the model on the documents
B) Splitting the documents into passages
C) Deleting every duplicate document
D) Translating the documents into English

Q6. State what a vector store holds in a retrieval pipeline.

A) The full text of every user question
B) The parameters of the language model
C) The rules of the organisation's policy
D) Numeric embeddings of the text passages

Q7. Name the Responsible Artificial Intelligence principle that requires personal data to be protected.

A) Privacy
B) Transparency
C) Accountability
D) Fairness

Q8. Identify the Responsible Artificial Intelligence principle requiring that a named person be answerable for a decision.

A) Fairness
B) Privacy
C) Accountability
D) Transparency

Q9. State the term for an attack in which text hidden inside an input changes the instructions a model follows.

A) Data poisoning
B) Model inversion
C) Prompt injection
D) Denial of service

Q10. Name the kind of explanation that accounts for one particular prediction.

A) A local explanation
B) A global explanation
C) A statistical summary
D) A model card

Q11. Identify the term for a synthetic image or video made to depict a real person convincingly.

A) A hallucination
B) An embedding
C) A guardrail
D) A deepfake

Q12. State the term for data that identifies a particular individual.

A) Aggregated statistical data
B) Personally identifiable information
C) Synthetic training data
D) Publicly licensed data

Q13. Name the practice of collecting only the data actually required for a stated purpose.

A) Data augmentation
B) Data replication
C) Data labelling
D) Data minimisation

Q14. Identify the setting that makes the output of a language model more predictable and less varied.

A) A high temperature
B) A larger context window
C) A low temperature
D) A longer system prompt

Q15. State what Artificial Intelligence governance provides within an organisation.

A) The hardware on which the models are trained
B) Policies, approval and oversight for the use of AI
C) The training data used to build the models
D) The programming language used for the code

Q16. Name the generative application that reduces a long document to its principal points.

A) Summarisation
B) Classification
C) Object detection
D) Clustering

Q17. Explain the difference between a generative model and a discriminative one.

A) One separates inputs, the other stores them
B) One composes new content, the other separates inputs
C) Both of the two compose entirely new content
D) Both of the two only separate inputs into classes

Q18. Distinguish a model that is interpretable by design from one explained after the event.

A) The first is readable in itself, the second is approximated
B) The second is readable in itself, the first approximated
C) Both of the two are readable in themselves
D) Both of the two are approximated after training

Q19. Describe what splitting documents into passages contributes to a retrieval pipeline.

A) It removes every factual error from the documents
B) It reduces the number of parameters in the model
C) It translates the documents into a common language
D) It lets only the relevant part be retrieved and supplied

Q20. Interpret a system that answers only from passages it has fetched and names the passage it used.

A) It is a model that has been fine-tuned on the data
B) It is a search engine with no model at all
C) It is a grounded, retrieval-augmented assistant
D) It is a rule-based expert system with citations

Q21. Predict what has to happen to a retrieval assistant when one of its source documents is revised.

A) The revised passage must be re-embedded and re-indexed
B) The whole language model must be retrained
C) Nothing at all, since the model adapts by itself
D) The system prompt has to be rewritten fully

Q22. Recognize the risk described: a web page an assistant is asked to read contains hidden text instructing it to disregard its own rules.

A) Hallucination of a source
B) Loss of the context window
C) Prompt injection
D) Overfitting to the page

Q23. Explain why asking a model to set out its intermediate steps can improve a multi-step answer.

A) The model consults an external calculator
B) The steps become part of the context it reasons from
C) The model is switched to a larger version
D) The model retrieves the answer from a database

Q24. Summarize the sources from which bias enters a deployed model.

A) The programming language and the hardware
B) The size of the context window in use
C) The number of users the system serves
D) The training data, the sampling and the labelling

Q25. Distinguish fairness from accuracy as objectives for a deployed model.

A) Fairness is overall correctness, accuracy concerns groups
B) Both of the two measure overall correctness
C) Accuracy is overall correctness, fairness concerns groups
D) Both of the two measure performance by group

Q26. Explain how a model can be accurate overall and still unfair to a particular group.

A) Accuracy is always computed on one group only
B) Fairness requires the accuracy to be exactly zero
C) The overall figure cannot be computed by group
D) A small group's errors barely affect the overall figure

Q27. Interpret an explanation that lists which features mattered most across the whole population of cases.

A) It is a global explanation
B) It is a local explanation
C) It is a fairness audit
D) It is a privacy assessment

Q28. Predict the effect on a language model's output of raising its temperature setting.

A) The output becomes shorter and more factual
B) The output becomes more varied and less predictable
C) The output becomes identical on every run
D) The output stops depending on the prompt

Q29. Describe what a guardrail on a deployed assistant does.

A) It increases the speed at which answers appear
B) It expands the amount of text the model can read
C) It retrains the model whenever an error occurs
D) It blocks or redirects inputs and outputs outside policy

Q30. Explain why removing names from a dataset does not always remove the privacy risk.

A) The names are still stored inside the model file
B) Remaining fields in combination can still identify a person
C) Anonymised data cannot be used to train a model
D) The risk transfers entirely to the model's output

Q31. Distinguish the training data of a model from the prompt it is given at the time of use.

A) A prompt shapes the model, training data one reply
B) Both of the two permanently change the model
C) Training data shapes the model, a prompt shapes one reply
D) Both of the two affect only a single reply

Q32. Explain why a generative model should not be the only source of a factual claim.

A) It is trained to produce fluent text, not verified fact
B) It cannot produce any factual statement at all
C) It only repeats sentences from its training data
D) It refuses to answer any factual question

Q33. Recognize the mitigation described: the assistant declines to answer and passes the question to a person when retrieval returns nothing relevant.

A) A refusal and hand-off path
B) A raised temperature setting
C) A larger context window
D) A fine-tuning of the model

Q34. Summarize the societal implications of Artificial Intelligence named on the syllabus.

A) Effects on the price of computing hardware only
B) Effects on the speed of internet connections only
C) Effects on the size of software companies only
D) Effects on employment, misinformation and equal treatment

Q35. Describe what the constraint section of a prompt contributes.

A) It supplies the documents to be searched

B) It bounds the length, scope and style of the reply

C) It sets the number of parameters used

D) It selects the hardware the model runs on

Q36. Explain why the output of a generative model should be labelled as machine generated.

A) The label improves the accuracy of the output

B) The label is required by the model's licence

C) A reader is entitled to know how the content arose

D) The label reduces the cost of each query

Q37. Distinguish Retrieval-Augmented Generation from simply training the model on more data.

A) Retrieval permanently alters the model's weights

B) Both of the two permanently alter the weights

C) Retrieval fetches sources at answer time without retraining

D) Neither of the two changes what the model can cite

Q38. Predict the difficulty of auditing a model whose internal reasoning cannot be inspected.

A) The model will refuse to produce any output

B) A prohibited factor may be in use without being visible

C) The model cannot be deployed on any hardware

D) The model becomes accurate but very slow

Q39. Explain what the transparency principle requires of an organisation deploying an Artificial Intelligence system.

A) That the source code is published without charge

B) That the training data is sold to other firms

C) That the model is retrained every single month

D) That people are told the system is in use and how it decides

Q40. Interpret the finding that a hiring model scores candidates from one university consistently lower than equally qualified candidates elsewhere.

A) The model has learned a bias present in past decisions

B) The model has been trained for too many epochs

C) The model's context window is set too small

D) The model's temperature has been set too high

----- End of Sample Question Set -----